

AIS Glass POC — Client Data Reconciliation

Oct 1, 2026 · @vijji babu

Does the proof of concept reconcile against the data Asahi India Glass supplied? Scope under test: 2 locations and 136 SKUs.

Verdict

The POC reconciles against the client data. The five source files are byte-identical to what was supplied, all 136 SKUs and both branches exist in the client's own order records, and the volume figures match exactly: 146,814 units ordered and 129,559 units despatched, over 70 weeks from 2025-W14 to 2026-W31.

Two items need attention before this is presented as complete:

1. **Two of the 136 SKUs produce no forecast.** They trained successfully but no champion model was selected for them, so the application reports 136 SKUs in scope while only 134 are forecastable.
2. **Order value differs by 0.48%** — ₹59.95Cr computed from the client file against ₹60.24Cr shown in the application. The cause is not yet established.

Source files are unaltered

Every file in the client delivery `OneDrive_1_08-09-2026.zip` is byte-for-byte identical to the copy the POC was built on. Compared by SHA-256 of the full file contents, not by name or size.

File	Size	Identical
Location Master.csv	0.02 MB	Yes
Orders & Receipts (Lead Time).xlsx	67.57 MB	Yes
Sales Data FY 24~26.xlsb	82.43 MB	Yes
Stock in Hand as on 1st Aug'26.xlsx	7.00 MB	Yes
Substitution Mapping.xlsx	0.22 MB	Yes

No client file was edited, reformatted or regenerated at any point in the build.

The scope exists in the client data

Both locations and all 136 SKUs were checked against the client's own Orders & Receipts file, joined on **Oracle No** — the order file's master key.

Check	Result
Branch BENGALURU present	Yes
Branch DELHI-1 present	Yes
SKUs configured in the POC	136
Of those, found in the client order file for these branches	136 — none missing
SKUs these two branches actually order	1,736

The 136 are a genuine subset of what the two branches buy, not a list containing anything the client did not supply.

Figures reconcile, except the value

The application's Overall Analysis page was recomputed directly from the client file over the same window and scope.

Measure	From client file	Application	Match
Order lines	30,217	—	—
Units ordered	146,814	146.8K	Exact
Units despatched	129,559	129.6K	Exact
Week range	2025-W14 → 2026-W31	2025-W14 → 2026-W31	Exact
Weeks covered	70	70	Exact
Order value	₹59.95Cr	₹60.24Cr	0.48% apart

The value gap is not a valuation-method difference

Both ways of valuing the client file agree with each other and give the same answer:

- Stored **MRP Value** column: ₹59.95Cr
- Recomputed **Quantity × MRP Rate**: ₹59.95Cr

- Rows where the two disagree by more than ₹1: **0 of 30,217**
- Rows where **MRP Value** is blank: **0**

So the ₹0.29Cr difference comes from somewhere inside the application's own panel, not from how the client file is read. **The cause has not been established and should not be presented as understood.**

What the scope covers

Measured on the client's Orders & Receipts file across all 775,912 lines.

Slice	Order lines	Units	Value	Share of network units
All 53 branches, all SKUs	775,912	2,602,392	₹857.12Cr	100%
The 2 POC branches, all SKUs	69,353	238,142	₹85.05Cr	9.2%
The 2 POC branches, 136 POC SKUs	30,217	146,814	₹59.95Cr	5.6%

Two readings matter here, and they are different claims.

Across the network, the POC covers about **9% of ordered units** and **7% of ordered value**. It is a pilot slice, and no figure in the application should be read as a national total — the workspace banner states this on every page.

Within the two branches, the 136 SKUs carry **61.6% of units** and **70.5% of value**, drawn from the 1,736 SKUs those branches order. The selection concentrates on high-value lines rather than sampling at random, so results are not representative of the long tail of 1,600 slower-moving SKUs.

Defect: two SKUs have no forecast

The application reports **136 SKUs in scope**, but only **134 have forecasts**. Two are silently absent.

SKU	Order lines in client data	Units	Model runs completed	Champion selected	Forecast rows
FG.T66.LFH.GCG212000W	111	526	26	No	0
FG.VWP.LFH.GCG212000W	53	182	26	No	0

Both have real demand in the client file. Both trained successfully — 26 completed model runs each, 8 models ineligible, which is normal. The break is at **champion selection**: no model was chosen for either scope, so no forecast could be generated.

This is a selection gap, not a data problem and not a training failure. The two SKUs are 0.5% of units in scope, so the headline figures are materially unaffected — but the count on screen says 136 when the forecastable number is 134, and that discrepancy should be closed rather than explained away.

Open items

- ☐ **Re-run champion selection across all scope kinds** so the two unforecast SKUs get a model. Until then, state the scope as 134 forecastable of 136 rather than 136.
- ☐ **Trace the ₹0.29Cr value difference** between the panel and the client file. Both are small, but an unexplained gap in a client-facing total is worth closing before it is asked about.
- ☐ **Decide whether to widen the SKU set.** 136 of 1,736 covers 70.5% of these branches' value; extending toward the tail would test the models on intermittent demand, which is where forecasting is hardest.

Nothing here blocks the POC being shown. The volume reconciliation is exact and the source data is provably untouched.

How each figure was produced

Every number above is reproducible from the client delivery and the running application. No figure is estimated.

File integrity. SHA-256 over the full contents of each file inside `OneDrive_1_08-09-2026.zip`, compared against `data/source/`. Hash equality, not size or filename.

Scope and volumes. Read directly from `Orders & Receipts (Lead Time).xlsx` — all 775,912 rows. Branches canonicalised by trim and uppercase; SKUs joined on `Oracle No`, which is the order file's master key, rather than `Material Code`. Filtered to the two branches and 136 SKUs, then to ISO weeks 2025-W14 through 2026-W31 to match the window the page reports.

Value. Computed twice and cross-checked: the stored `MRP Value` column, and `Quantity × MRP Rate` recomputed per line.

Forecast coverage and the two missing SKUs. Queried from the application's own database — distinct SKUs on `forecast_row`, champion rows on `champion_selection`, and model outcomes on `model_run` grouped by status.

Scope under test. Branches and SKUs taken from the deployment's own configuration, cross-checked against the deployed bundle manifest, which records `{branches: [BENGALURU, DELHI-1], sku_count: 136}`.